

PS2-EDA

September 26, 2023

0.1 Problem set 2

0.1.1 Problem 0

-2 points for every missing green OK sign. If you don't run the cell below, that's -14 points.

Make sure you are in the DATA1030 environment.

```
[ ]: from __future__ import print_function
from packaging.version import parse as Version
from platform import python_version

OK = '\x1b[42m[ OK ]\x1b[0m'
FAIL = "\x1b[41m[FAIL]\x1b[0m"

try:
    import importlib
except ImportError:
    print(FAIL, "Python version 3.10 is required,"
              " but %s is installed." % sys.version)

def import_version(pkg, min_ver, fail_msg=""):
    mod = None
    try:
        mod = importlib.import_module(pkg)
        if pkg in {'PIL'}:
            ver = mod.VERSION
        else:
            ver = mod.__version__
        if Version(ver) == Version(min_ver):
            print(OK, "%s version %s is installed."
                  % (lib, min_ver))
        else:
            print(FAIL, "%s version %s is required, but %s installed."
                  % (lib, min_ver, ver))
    except ImportError:
        print(FAIL, '%s not installed. %s' % (pkg, fail_msg))
    return mod
```

```

# first check the python version
pyversion = Version(python_version())

if pyversion >= Version("3.11.4"):
    print(OK, "Python version is %s" % pyversion)
elif pyversion < Version("3.11"):
    print(FAIL, "Python version 3.11 is required,"
              " but %s is installed." % pyversion)
else:
    print(FAIL, "Unknown Python version: %s" % pyversion)

print()
requirements = {'numpy': "1.24.4", 'matplotlib': "3.7.2", 'sklearn': "1.3.0",
                'pandas': "2.0.3", 'xgboost': "1.7.6", 'shap': "0.42.1", ↴
                'seaborn': "0.12.2"}

# now the dependencies
for lib, required_version in list(requirements.items()):
    import_version(lib, required_version)

```

[OK] Python version is 3.11.4

[OK] numpy version 1.24.4 is installed.
[OK] matplotlib version 3.7.2 is installed.
[OK] sklearn version 1.3.0 is installed.
[OK] pandas version 2.0.3 is installed.
[OK] xgboost version 1.7.6 is installed.
[OK] shap version 0.42.1 is installed.
[OK] seaborn version 0.12.2 is installed.

0.1.2 Problem 1

One of the datasets we will be working with this semester is the kaggle house price dataset. The goal of this problem set is to use this dataset to practice dataframe manipulations and perform EDA.

Carefully read the description of the dataset which is located in the `data` folder. This is a very well documented dataset. Whenever you work with a dataset, it is highly recommended that you prepare a similar description if it is not readily available. Specific things to note:

- each feature is described in full detail,
- the meaning of continuous features is explained and their unit is provided (e.g., lot size is measured in square feet),
- each category in a categorical or ordinal feature is spelled out and explained.

0.1.3 Problem 1a: Read in the data (5 points)

The kaggle house price dataset is located in the `data` folder.

First, read the data into a pandas data frame and display the data frame below. You might encounter error messages and other issues along the way. Please diagnose and resolve them.

```
[ ]: # your code here
```

```
[ ]: import pandas as pd
```

```
[ ]: pip install openpyxl
```

```
Collecting openpyxl
  Using cached openpyxl-3.1.2-py2.py3-none-any.whl (249 kB)
Requirement already satisfied: et-xmlfile in
/Users/apple/opt/anaconda3/envs/data1030/lib/python3.11/site-packages (from
openpyxl) (1.1.0)
Installing collected packages: openpyxl
Successfully installed openpyxl-3.1.2
Note: you may need to restart the kernel to use updated packages.
```

```
[ ]: df = pd.read_excel("/Users/apple/Desktop/Data 1030/
˓→github-classroom-Data1030-Xiner Zhao/ps2-eda-XXXXiner/data/train.xlsx", u
˓→sheet_name="data")
df.head()
```

```
[ ]:   MSSubClass MSZoning LotFrontage LotArea Street Alley LotShape \
```

```
0       60      RL     65.0    8450    Pave    NaN    Reg
1       20      RL     80.0    9600    Pave    NaN    Reg
2       60      RL     68.0   11250    Pave    NaN    IR1
3       70      RL     60.0    9550    Pave    NaN    IR1
4       60      RL     84.0   14260    Pave    NaN    IR1
```

```
  LandContour Utilities LotConfig ... PoolArea PoolQC Fence MiscFeature \
```

```
0      Lvl    AllPub    Inside ...        0    NaN    NaN    NaN
1      Lvl    AllPub        FR2 ...        0    NaN    NaN    NaN
2      Lvl    AllPub    Inside ...        0    NaN    NaN    NaN
3      Lvl    AllPub    Corner ...        0    NaN    NaN    NaN
4      Lvl    AllPub        FR2 ...        0    NaN    NaN    NaN
```

```
  MiscVal MoSold YrSold SaleType SaleCondition SalePrice
```

```
0      0      2    2008      WD      Normal  208500
1      0      5    2007      WD      Normal  181500
2      0      9    2008      WD      Normal  223500
3      0      2    2006      WD  Abnorml  140000
4      0     12    2008      WD      Normal  250000
```

```
[5 rows x 80 columns]
```

0.1.4 Problem 1b: dataframe filtering

Answer the following data filtering related questions.

Please make it easy for the TAs to grade your solution so print no more and no less than the necessary info. For example, if we ask how many columns there are in the dataset, print out that one number instead of the full shape of the data frame.

Q1 When you buy a house, a surprising amount of time is spent assessing the basement. List the names of the basement-related features! How many basement-related features are there? (2 points)

```
[ ]: # your code here
```

```
columns = df.columns
```

```
print(columns)
```

```
Index(['MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street', 'Alley',
       'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope',
       'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle',
       'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'RoofStyle',
       'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'MasVnrArea',
       'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond',
       'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2',
       'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating', 'HeatingQC',
       'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF',
       'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath',
       'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual', 'TotRmsAbvGrd',
       'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType', 'GarageYrBlt',
       'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond',
       'PavedDrive', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch',
       'ScreenPorch', 'PoolArea', 'PoolQC', 'Fence', 'MiscFeature', 'MiscVal',
       'MoSold', 'YrSold', 'SaleType', 'SaleCondition', 'SalePrice'],
      dtype='object')
```

After we print all the name of columns in dataframe, we found that the keyword of basement-related feature is 'Bsmt', so we need to find out which columns contain the key word 'Bsmt'. Then we have the code below:

```
[ ]: keyword = 'Bsmt'
```

```
selected_columns = [col for col in columns if keyword in col]
```

```
print('The names of the basement-related features are: \n', selected_columns)
```

```
print('The quantity of basement-related features is:', len(selected_columns))
```

The names of the basement-related features are:

```
['BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1',
 'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'BsmtFullBath',
 'BsmtHalfBath']
```

The quantity of basement-related features is: 11

Q2 How many houses have an excellent overall quality? (1 point)

```
[ ]: # your code here  
print(df['OverallQual'].head())
```

```
0    7  
1    6  
2    7  
3    7  
4    8  
Name: OverallQual, dtype: int64
```

We can find the variable ‘OverallQual’ is in integer format and it is not in a string format, so we need to check the description of the dataset and find which number means ‘Excellent’. We can find that 9 means excellent from data description. Then we have the code below:

```
[ ]: df_OverallQual_Excellent = df[df['OverallQual'] == 9]  
print('The quantity of houses which have an excellent overall quality is:  
    ↵',df_OverallQual_Excellent.shape[0])
```

```
The quantity of houses which have an excellent overall quality is: 43
```

Q3 How many houses were built on or before the year 2000? (1 point)

```
[ ]: # your code here  
print(df['YearBuilt'].head())
```

```
0    2003  
1    1976  
2    2001  
3    1915  
4    2000  
Name: YearBuilt, dtype: int64
```

```
[ ]: df_YearBuilt_2000 = df[df['YearBuilt'] <= 2000]  
print('The quantity of houses which were built on or before the year 2000 is:  
    ↵',df_YearBuilt_2000.shape[0])
```

```
The quantity of houses which were built on or before the year 2000 is: 1096
```

Q4 How many houses have a pool and central airconditioning? (2 points)

```
[ ]: # your code here  
print(df['CentralAir'].head())  
print(df['PoolArea'].head())
```

```
0    Y  
1    Y  
2    Y  
3    Y  
4    Y  
Name: CentralAir, dtype: object  
0    0
```

```
1    0  
2    0  
3    0  
4    0  
Name: PoolArea, dtype: int64
```

```
[ ]: df_Pool_AirCon = df[(df['CentralAir'] == 'Y') & (df['PoolArea'] > 0)]  
print('The quantity of houses which have a pool and central airconditioning is:  
', df_Pool_AirCon.shape[0])
```

The quantity of houses which have a pool and central airconditioning is: 7

0.1.5 Problem 2a: EDA

Answer the following EDA related questions.

The sequence of questions here are typical things to ask when you perform EDA on a new dataset. First you always want to know how many data points and features you have, and whether they are continuous or categorical. You should then take a closer look at the target variable. We will study the properties of the features and the relationships between the features and the target variable in 2b.

Q1 How many rows and columns do we have in the dataframe? (2 point)

```
[ ]: # your code here  
print('The quantity of row in dataframe is:', df.shape[0])  
print('The quantity of column in dataframe is:', df.shape[1])
```

The quantity of row in dataframe is: 1460

The quantity of column in dataframe is: 80

Q2 What are the data types of the columns? Make sure that the output is not truncated and you see the type of each column. (2 points)

```
[ ]: # your code here  
pd.set_option('display.max_columns', None) # In order to display all output  
pd.set_option('display.max_rows', None) # In order to display all output  
df.dtypes
```

```
[ ]: MSSubClass      int64  
MSZoning        object  
LotFrontage     float64  
LotArea         int64  
Street          object  
Alley           object  
LotShape         object  
LandContour     object  
Utilities        object  
LotConfig        object  
LandSlope        object
```

Neighborhood	object
Condition1	object
Condition2	object
BldgType	object
HouseStyle	object
OverallQual	int64
OverallCond	int64
YearBuilt	int64
YearRemodAdd	int64
RoofStyle	object
RoofMatl	object
Exterior1st	object
Exterior2nd	object
MasVnrType	object
MasVnrArea	float64
ExterQual	object
ExterCond	object
Foundation	object
BsmtQual	object
BsmtCond	object
BsmtExposure	object
BsmtFinType1	object
BsmtFinSF1	int64
BsmtFinType2	object
BsmtFinSF2	int64
BsmtUnfSF	int64
TotalBsmtSF	int64
Heating	object
HeatingQC	object
CentralAir	object
Electrical	object
1stFlrSF	int64
2ndFlrSF	int64
LowQualFinSF	int64
GrLivArea	int64
BsmtFullBath	int64
BsmtHalfBath	int64
FullBath	int64
HalfBath	int64
BedroomAbvGr	int64
KitchenAbvGr	int64
KitchenQual	object
TotRmsAbvGrd	int64
Functional	object
Fireplaces	int64
FireplaceQu	object
GarageType	object

```
GarageYrBlt      float64
GarageFinish     object
GarageCars       int64
GarageArea       int64
GarageQual       object
GarageCond       object
PavedDrive      object
WoodDeckSF      int64
OpenPorchSF     int64
EnclosedPorch   int64
3SsnPorch       int64
ScreenPorch     int64
PoolArea        int64
PoolQC          object
Fence           object
MiscFeature     object
MiscVal          int64
MoSold          int64
YrSold          int64
SaleType         object
SaleCondition   object
SalePrice        int64
dtype: object
```

```
[ ]: pd.reset_option("display.max_columns") # Restore the setting
pd.reset_option("display.max_rows") # Restore the setting
```

Q3 The ML target variable in this dataset is the sale price. We will develop ML pipelines to predict this variable based on the other features.

Is this column continuous or categorical? Please use .describe or .value_counts to take a quick look at this feature. (2 points)

```
[ ]: # your code here
df['SalePrice'].describe()
```


```
[ ]: count      1460.000000
mean      180921.195890
std       79442.502883
min       34900.000000
25%      129975.000000
50%      163000.000000
75%      214000.000000
max       755000.000000
Name: SalePrice, dtype: float64
```

From the output, we can find that the sale price is continuous.

Q4 Visualize the target variable. Don't forget the axis labels and graph title. Make sure to use

appropriate arguments to best display the data. (4 points)

```
[ ]: # your code here
import numpy as np
import matplotlib
from matplotlib import pylab as plt
plt.figure(figsize=(8,6))
df['SalePrice'].plot.hist(bins = int(np.sqrt(df.shape[0])))
plt.xlabel('SalePrice')
plt.ylabel('count')
plt.title('The SalePrice of House')
plt.show()
```



0.1.6 Problem 2b: visualization (15 points)

Find one continuous, one ordinal, and one categorical feature that strongly correlates with the sale price. Create figures that illustrate your selected features and the sale price.


Don't forget to add axis labels and titles, and find appropriate arguments. Write figure captions to explain what the figure shows.

We will cover techniques in a week or two to quantitatively assess how strongly each feature correlates with the target variable. For now, we do a qualitative/visual assessment.

```
[ ]: # your code here
# (1) Continuous Feature 'GrLivArea' VS continuous feature 'SalePrice'
    ↵(GrLivArea: Above grade (ground) living area square feet)
df['GrLivArea'].describe()
```

```
[ ]: count      1460.000000
mean       1515.463699
std        525.480383
min       334.000000
25%      1129.500000
50%      1464.000000
75%      1776.750000
max      5642.000000
Name: GrLivArea, dtype: float64
```

```
[ ]: df.plot.scatter('GrLivArea','SalePrice',figsize=(8,6),alpha=0.2,s=20)
plt.xlabel('GrLivArea: Above ground living area square feet')
plt.ylabel('SalePrice')
plt.title('SalePrice in different GrLivArea')
plt.show()
```




The figure shows how the sale price of a house changes when its above ground living area changes.

```
[ ]: # (2) Ordinal Feature 'GarageFinish' VS continuous feature 'SalePrice'  
    ↴(GarageFinish: Interior finish of the garage)  
df['GarageFinish'].value_counts()
```

```
[ ]: GarageFinish  
Unf      605  
RFn     422  
Fin      352  
Name: count, dtype: int64
```

```
[ ]: df[['SalePrice','GarageFinish']].boxplot(by='GarageFinish',figsize=(8,6))  
plt.ylabel('SalePrice')  
plt.xticks([1,2,3],['Finished','Rough Finished','Unfinished'])  
plt.title('SalePrice in different GarageFinish')  
plt.show()
```




The figure shows how the sale price of a house changes when its garage is finished, rough finished

or unfinished.

```
[ ]: # (3) Categorical feature 'GarageType' VS continuous feature 'SalePrice'  
    ↪(GarageType: Garage location)  
df['GarageType'].value_counts()
```

```
[ ]: GarageType  
Attchd      870  
Detchd      387  
BuiltIn      88  
Basment     19  
CarPort      9  
2Types       6  
Name: count, dtype: int64
```

```
[ ]: dataset = [df[df['GarageType']=='2Types']['SalePrice'].values,  
               df[df['GarageType']=='Attchd']['SalePrice'].values,  
               df[df['GarageType']=='Basment']['SalePrice'].values,  
               df[df['GarageType']=='BuiltIn']['SalePrice'].values,  
               df[df['GarageType']=='CarPort']['SalePrice'].values,  
               df[df['GarageType']=='Detchd']['SalePrice'].values]  
  
plt.figure(figsize=(10,8))  
plt.violinplot(dataset = dataset)  
plt.  
    ↪xticks([1,2,3,4,5,6],['2Types','Attchd','Basment','BuiltIn','CarPort','Detchd'])  
plt.ylabel('SalePrice')  
plt.legend(['2Types:More than one type of garage','Attchd:Attached to  
    ↪home','Basment:Basement Garage','BuiltIn:Garage part of house','CarPort;Car  
    ↪Port','Detchd:Detached from home'])  
plt.title('SalePrice in different GarageType')  
plt.show()
```



The figure shows how the sale price of a house changes when its type of garage changes.