

# PS3\_Xiner Zhao

October 8, 2023

## 0.1 Problem set 3

### 0.1.1 Problem 0

Run the cell below to make sure you are in the data1030 coding environment.

We will deduct 2 points for every missing OK sign. (If you don't run the cell, that's -14 points.)

```
[ ]: from __future__ import print_function
from packaging.version import parse as Version
from platform import python_version

OK = '\x1b[42m[ OK ]\x1b[0m'
FAIL = "\x1b[41m[FAIL]\x1b[0m"

try:
    import importlib
except ImportError:
    print(FAIL, "Python version 3.11 is required,"
              " but %s is installed." % sys.version)

def import_version(pkg, min_ver, fail_msg=""):
    mod = None
    try:
        mod = importlib.import_module(pkg)
        if pkg in {'PIL'}:
            ver = mod.VERSION
        else:
            ver = mod.__version__
        if Version(ver) == Version(min_ver):
            print(OK, "%s version %s is installed."
                  % (lib, min_ver))
        else:
            print(FAIL, "%s version %s is required, but %s installed."
                  % (lib, min_ver, ver))
    except ImportError:
        print(FAIL, '%s not installed. %s' % (pkg, fail_msg))
    return mod
```

```

# first check the python version
pyversion = Version(python_version())

if pyversion >= Version("3.11.4"):
    print(OK, "Python version is %s" % pyversion)
elif pyversion < Version("3.11"):
    print(FAIL, "Python version 3.11 is required,"
              " but %s is installed." % pyversion)
else:
    print(FAIL, "Unknown Python version: %s" % pyversion)

print()
requirements = {'numpy': "1.24.4", 'matplotlib': "3.7.2", 'sklearn': "1.3.0",
                'pandas': "2.0.3", 'xgboost': "1.7.6", 'shap': "0.42.1", ↴
                'seaborn': "0.12.2"}

# now the dependencies
for lib, required_version in list(requirements.items()):
    import_version(lib, required_version)

```

[ OK ] Python version is 3.11.4

[ OK ] numpy version 1.24.4 is installed.  
[ OK ] matplotlib version 3.7.2 is installed.  
[ OK ] sklearn version 1.3.0 is installed.  
[ OK ] pandas version 2.0.3 is installed.  
[ OK ] xgboost version 1.7.6 is installed.  
[ OK ] shap version 0.42.1 is installed.  
[ OK ] seaborn version 0.12.2 is installed.

### 0.1.2 Problem 1a, diabetes dataset EDA (10 points)

We will work with the diabetes dataset in problem 1. Please carefully read the description [here](#) and [here](#). Read the data into a pandas dataframe using the txt file linked [here](#). Perform EDA by going through questions 1-4 from PS2 2a.

Additionally, prepare two visualizations using column pairs. The two figures should be different types but one of the columns for both plots should be the target variable. Include a caption for each of the figures describing what the plot shows.

[ ]: # your code here

[ ]: import numpy as np  
data = np.loadtxt("https://www4.stat.ncsu.edu/~boos/var.select/diabetes.tab.  
˓→txt", dtype = 'object', delimiter = '\t')  
print(data)

```

[['AGE' 'SEX' 'BMI' ... 'S5' 'S6' 'Y']
 ['59' '2' '32.1' ... '4.8598' '87' '151']
 ['48' '1' '21.6' ... '3.8918' '69' '75']
 ...
 ['60' '2' '24.9' ... '4.1271' '95' '132']
 ['36' '1' '30' ... '5.1299' '85' '220']
 ['36' '1' '19.6' ... '4.5951' '92' '57']]

```

```

[ ]: import pandas as pd
df = pd.DataFrame(data, columns = data[0])
df.drop([0], inplace=True)
df = pd.DataFrame(df,dtype=float)
df['AGE'] = df['AGE'].astype(int)
df['SEX'] = df['SEX'].astype(int)
print(df)

```

	AGE	SEX	BMI	BP	S1	S2	S3	S4	S5	S6	Y
1	59	2	32.1	101.00	157.0	93.2	38.0	4.00	4.8598	87.0	151.0
2	48	1	21.6	87.00	183.0	103.2	70.0	3.00	3.8918	69.0	75.0
3	72	2	30.5	93.00	156.0	93.6	41.0	4.00	4.6728	85.0	141.0
4	24	1	25.3	84.00	198.0	131.4	40.0	5.00	4.8903	89.0	206.0
5	50	1	23.0	101.00	192.0	125.4	52.0	4.00	4.2905	80.0	135.0
..	...	...	...	...	...	...	...	...	...	...	...
438	60	2	28.2	112.00	185.0	113.8	42.0	4.00	4.9836	93.0	178.0
439	47	2	24.9	75.00	225.0	166.0	42.0	5.00	4.4427	102.0	104.0
440	60	2	24.9	99.67	162.0	106.6	43.0	3.77	4.1271	95.0	132.0
441	36	1	30.0	95.00	201.0	125.2	42.0	4.79	5.1299	85.0	220.0
442	36	1	19.6	71.00	250.0	133.2	97.0	3.00	4.5951	92.0	57.0

[442 rows x 11 columns]

```

[ ]: print('The quantity of row in dataframe is:',df.shape[0])
print('The quantity of column in dataframe is:',df.shape[1])
print(df.dtypes)

```


The quantity of row in dataframe is: 442  
The quantity of column in dataframe is: 11

	AGE	SEX	BMI	BP	S1	S2	S3	S4	S5	S6	Y
AGE	int64										
SEX		int64									
BMI			float64								
BP				float64							
S1					float64						
S2						float64					
S3							float64				
S4								float64			
S5									float64		
S6										float64	
Y											float64
											object

```
[ ]: df['Y'].describe()
```

```
[ ]: count      442.000000
mean       152.133484
std        77.093005
min       25.000000
25%       87.000000
50%      140.500000
75%      211.500000
max      346.000000
Name: Y, dtype: float64
```


```
[ ]: import matplotlib
from matplotlib import pylab as plt
plt.figure(figsize=(8,6))
plt.hist(df['Y'], bins = int(np.sqrt(df.shape[0])))
plt.xlabel('Y : Disease Progression One Year After Baseline')
plt.ylabel('count')
plt.title('The Distribution of Y')
plt.show()
```



```
[ ]: df['S1'].describe()
```

```
[ ]: count      442.000000
mean       189.140271
std        34.608052
min        97.000000
25%       164.250000
50%       186.000000
75%       209.750000
max       301.000000
Name: S1, dtype: float64
```

```
[ ]: nbins = 40
heatmap, xedges, yedges = np.histogram2d(df['S1'], df['Y'], bins=nbins)
extent = [xedges[0], xedges[-1], yedges[0], yedges[-1]]
heatmap[heatmap == 0] = 0.1
plt.figure(figsize=(8,6))
plt.imshow(heatmap.T, origin='lower', vmin=0)
plt.xlabel('S1 : Total Serum Cholesterol')
plt.ylabel('Y : Disease Progression One Year After Baseline')
plt.xticks(np.arange(nbins+1)[::4], xedges[::4].astype(int))
plt.yticks(np.arange(nbins+1)[::4], yedges[::4].astype(int))
plt.colorbar(label='count')
plt.show()
```




From this heatmap, we can easily figure out that the value of quantitative measure of disease progression is larger if people's total serum cholesterol is between 158 and 219, and many people whose total serum cholesterol is around 180 have around 185 value in the quantitative measure of disease progression.

```
[ ]: df['SEX'].value_counts()
```

```
[ ]: SEX
1    235
2    207
Name: count, dtype: int64
```

```
[ ]: dataset = [df[df['SEX']==1]['Y'].values,
               df[df['SEX']==2]['Y'].values]
plt.figure(figsize=(8,6))
plt.violinplot(dataset)
plt.xticks([1,2])
plt.ylabel('Y : Disease Progression One Year After Baseline')
```

```
plt.title('Disease progression in difference sex')
plt.show()
```



From this violinplot, we can easily figure out that the quantitative measure of disease progression (Y) of people, whose sex equals to 1, ranges wider than that of people whose sex equals to 2, and also the average of Y of people, whose sex equals to 1, is lower than that of people whose sex equals to 2.

### 0.1.3 Problem 1b, basic split (10 points)

Write a general function that performs basic splitting on a dataset, while also conducting integrity tests on both its inputs and outputs. Let's call the function `basic_split`. It takes the following arguments as inputs: feature matrix (`X`), a target variable (`y`), `train_size`, `val_size`, `test_size`, and `random_state`. The output of the function should be: `X_train`, `y_train`, `X_val`, `y_val`, `X_test`, `y_test`.

Perform the following tests inside the function and raise a value error with a message if one or more of the tests fail.

Test the inputs: - the sum of `train_size`, `val_size`, `test_size` is 1 - `random_state` is an integer

Use `train_test_split`.

Test the outputs: - check that the output sizes are what you want them to be

Note that in principle you could add more tests. For example: - test if X is a 2d pandas data frame  
- test if y is a 1d pandas series - check if the number of rows in X is the same as the length of y

These tests are included in train\_test\_split, so we don't need to explicitly add them to our code. However, if you were to split your data manually without train\_test\_split, it would be a good idea to add these additional tests and potentially more.

Apply the function to the diabetes dataset with train\_size = 0.6, val\_size = 0.2, and test\_size = 0.2. Print out the head of X\_train, X\_val, and X\_test. Make sure that you get the same points in each set every time you rerun the cell (a.k.a., check for reproducability).

This function is general purpose, you'll be able to reuse it for any project if you want to perform basic split on your data.

```
[ ]: # your code here

from sklearn.model_selection import train_test_split
def basic_split(X,y,train_size,val_size,test_size,random_state):

    if train_size+val_size+test_size != 1:
        print('The sum of train_size, val_size and test_size is not equal to 1! Please modify your split size.')
    elif type(random_state) != int:
        print('The random_state is not an integer! Please input an integer for random_state.')
    elif X.ndim != 2:
        print('The dimension of X is not 2! Please input a 2D pandas dataframe for X.')
    elif y.ndim != 1:
        print('The dimension of y is not 1! Please input a 1D pandas series for y.')
    elif X.shape[0] != len(y):
        print('The number of rows in X is not same as the length of y! Please check your input data set X and y.')
    else:
        # First split to separate out the training set:
        X_train, X_other, y_train, y_other = train_test_split(X,y,train_size = train_size,random_state = random_state)
        # Second split to separate out the validation set and the test set:
        X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_size=val_size/(val_size+test_size),random_state=random_state)
        if round(X_train.shape[0]/X.shape[0],1) != train_size or round(len(y_train)/len(y),1) != train_size:
            print('The number of rows in X_train or y_train is not equal to the train_size you input! Please check.')
```

```

        X_train,y_train,X_val,y_val,X_test,y_test = ↵
↳None,None,None,None,None
            return X_train,y_train,X_val,y_val,X_test,y_test
        elif round(X_val.shape[0]/X.shape[0],1) != val_size or round(len(y_val)/
↳len(y),1) != val_size:
            print('The number of rows in X_val or y_val is not equal to the
↳val_size you input! Please check.')
        X_train,y_train,X_val,y_val,X_test,y_test = ↵
↳None,None,None,None,None
            return X_train,y_train,X_val,y_val,X_test,y_test
        elif round(X_test.shape[0]/X.shape[0],1) != test_size or
↳round(len(y_test)/len(y),1) != test_size:
            print('The number of rows in X_test or y_test is not equal to the
↳test_size you input! Please check.')
        X_train,y_train,X_val,y_val,X_test,y_test = ↵
↳None,None,None,None,None
            return X_train,y_train,X_val,y_val,X_test,y_test
    else:
        return X_train,y_train,X_val,y_val,X_test,y_test

```

```
[ ]: X_columns = df.columns != 'Y'
X_df = df.loc[:,X_columns]
y_df = df['Y']
print(X_df)
print(y_df)
```

	AGE	SEX	BMI	BP	S1	S2	S3	S4	S5	S6
1	59	2	32.1	101.00	157.0	93.2	38.0	4.00	4.8598	87.0
2	48	1	21.6	87.00	183.0	103.2	70.0	3.00	3.8918	69.0
3	72	2	30.5	93.00	156.0	93.6	41.0	4.00	4.6728	85.0
4	24	1	25.3	84.00	198.0	131.4	40.0	5.00	4.8903	89.0
5	50	1	23.0	101.00	192.0	125.4	52.0	4.00	4.2905	80.0
..	...	...	...	...	...	...	...	...	...	...
438	60	2	28.2	112.00	185.0	113.8	42.0	4.00	4.9836	93.0
439	47	2	24.9	75.00	225.0	166.0	42.0	5.00	4.4427	102.0
440	60	2	24.9	99.67	162.0	106.6	43.0	3.77	4.1271	95.0
441	36	1	30.0	95.00	201.0	125.2	42.0	4.79	5.1299	85.0
442	36	1	19.6	71.00	250.0	133.2	97.0	3.00	4.5951	92.0

[442 rows x 10 columns]

1	151.0
2	75.0
3	141.0
4	206.0
5	135.0
..	..
438	178.0

```

439    104.0
440    132.0
441    220.0
442     57.0
Name: Y, Length: 442, dtype: float64

```

```
[ ]: random_state_df = 1030
train_size_df = 0.6
val_size_df = 0.2
test_size_df = 0.2
X_train_df,y_train_df,X_val_df,y_val_df,X_test_df,y_test_df =
    basic_split(X_df,y_df,train_size_df,val_size_df,test_size_df,random_state_df)
```

```
[ ]: print(X_train_df.head())
print(X_train_df.shape)

print(X_val_df.head())
print(X_val_df.shape)

print(X_test_df.head())
print(X_test_df.shape)
```

	AGE	SEX	BMI	BP	S1	S2	S3	S4	S5	S6
420	43	1	21.3	79.0	141.0	78.8	53.0	3.0	3.8286	90.0
246	41	1	23.1	86.0	148.0	78.0	58.0	3.0	4.0943	60.0
76	46	2	23.5	87.0	181.0	114.8	44.0	4.0	4.7095	98.0
382	29	2	18.1	73.0	158.0	99.0	41.0	4.0	4.4998	78.0
240	55	1	28.2	91.0	250.0	140.2	67.0	4.0	5.3660	103.0
(265,	10)									
41	50	2	25.6	101.0	229.0	162.2	43.0	5.0	4.7791	114.0
47	33	1	25.3	85.0	155.0	85.0	51.0	3.0	4.5539	70.0
185	53	1	28.6	88.0	171.0	98.8	41.0	4.0	5.0499	99.0
370	46	1	29.9	83.0	171.0	113.0	38.0	4.5	4.5850	98.0
107	22	1	19.3	82.0	156.0	93.2	52.0	3.0	3.9890	71.0
(88,	10)									
407	33	1	18.9	70.0	162.0	91.8	59.0	3.0	4.0254	58.0
221	55	2	22.7	93.0	154.0	94.2	53.0	3.0	3.5264	75.0
405	44	1	31.4	115.0	165.0	97.6	52.0	3.0	4.3438	89.0
164	53	2	33.1	117.0	183.0	119.0	48.0	4.0	4.3820	106.0
52	65	2	27.9	103.0	159.0	96.8	42.0	4.0	4.6151	86.0
(89,	10)									

**Q1b Check:** Is your function reproducible? How do you know? (1 point)

```
[ ]: result = dict()
for I in (1,2,3,4):
```

```

    result[I] =_
↪basic_split(X_df,y_df,train_size_df,val_size_df,test_size_df,random_state_df)[-1]
    print('The %sth result is:\n '%I,result[I])
#print(result)

```

The 1th result is:

407	72.0
221	78.0
405	293.0
164	131.0
52	225.0
...	
116	229.0
203	196.0
343	178.0
65	71.0
212	70.0

Name: Y, Length: 89, dtype: float64

The 2th result is:

407	72.0
221	78.0
405	293.0
164	131.0
52	225.0
...	
116	229.0
203	196.0
343	178.0
65	71.0
212	70.0

Name: Y, Length: 89, dtype: float64

The 3th result is:

407	72.0
221	78.0
405	293.0
164	131.0
52	225.0
...	
116	229.0
203	196.0
343	178.0
65	71.0
212	70.0

Name: Y, Length: 89, dtype: float64

The 4th result is:

407	72.0
221	78.0
405	293.0

```

164    131.0
52     225.0
...
116    229.0
203    196.0
343    178.0
65     71.0
212    70.0
Name: Y, Length: 89, dtype: float64

```

From the output, we can find that no matter how many times we rerun the code, the result will be always same if we set up a fixed random\_state.

#### 0.1.4 Problem 1c, stratified regression (5 points)

I mentioned in class that sklearn's splitting methods can only stratify with respect to a classification target variable. However you might encounter scenarios where it is necessary to stratify on a regression target variable. For example, it is a good idea to stratify if the distribution of the regression target variable is heavy-tailed (e.g., exponential, log-normal). If you do not stratify, some of your sets might not contain rare values from the heavy tail thus throwing off the regression model.

Either come up with an algorithm to stratify with respect to a regression target variable and write pseudocode in a markdown cell below to explain the steps of the alrogithm; or do some reading online and write a paragraph or two to explain already existing algorithms/approaches to solve the problem and add your references.

[ ]: # your answer here

Suppose we have a data set DF with a regression target variable DF\_y and other features (DF\_X1,DF\_X2,...,DF\_Xn): import numpy as np

```

X_columns = DF.columns != 'DF_y'
DF_X = DF.loc[:,X_columns]
DF_y = DF['DF_y']
nrows = DF.shape[0]
bins = np.linspace(np.min(DF['DF_y']),np.max(DF['DF_y']),int(np.sqrt(nrows)))
y_binned = np.digitize(DF['DF_y'],bins)
X_train, X_other, y_train, y_other = train_test_split(DF_X,DF_y,train_size=train_size,
stratify=y_binned, random_state=random_state)
y_binned_other = y_binned[y_other.index.tolist()]
X_val, X_test, y_val, y_test = train_test_split(X_other,y_other,train_size=val_size/(val_size+test_size),stratify=y_binned_other)

```

#### Explanation:

- (1) Firstly, extract the feature data set X and the regression target variable y as usual.

- (2) Secondly, set up a bin in an arithmetic progression, of which the start value is the minimum of  $y$  and the end value is the maximum of  $y$ , and the common difference is  $\text{int}(\text{np.sqrt}(\text{nrows}))$ .
- (3) Thirdly, split the regression target variable  $y$  into different interval(groups) according to the bin by `np.digitize`, and we can denote the group name each  $y$  belongs to as  $y\_binned$ . That is, for example, if  $y[1]$  belongs to the 10th interval(or the 10th group), the  $y\_binned[1]$  equals to 10, and same as other value in  $y$ . Then, for each value in  $y$ , they belong to an interval(group). So we can convert a continuous variable into a categorical variable.
- (4) Finally, use the `train_test_split` function to split the data set into train, validation, test data set in the size we want as usual.

**Reference:** [Michael J. Sanders](#)

### 0.1.5 Problem 2

We will work with the [hand postures dataset](#) in problem 2. Please carefully read the dataset description and perform as much EDA as you can on this dataset. The EDA is not graded but it will help you to correctly answer 2a and 2b.

This dataset has group structure: 14 users performing 5 different hand postures while wearing sensors attached to a left-handed glove. Two different ML questions can be asked using this dataset. We will explore how splitting differs for both questions in 2a and 2b.

**A note on splitting non-IID data:** The Independent and Identically Distributed (IID) assumption asserts two foundational premises: (1) the values of one datapoint (i.e. row) in our dataset have no impact on the values of another datapoint, and (2) all of the datapoints come from the same distribution. When this assumption holds, we can perform basic splitting techniques as explored above, since the make-up of our training, validation, and test sets will be independent of one another. However, many ML problems involve using data that does not conform to the IID assumption, such as data with a group or time-series structure. When your dataset has a group structure, datapoints from one group may be drawn from a different distribution than datapoints from another group. Furthermore, the datapoints in a time-series dataset have some correlation with one another, and therefore the values of one point may implicitly hint at the values of another. This raises the issue of data leakage; the model can train on data that unfairly gives it an advantage when predicting the values in its test set. For instance, if you are trying to predict the price of a stock, and the model trains on share price data from 2023, it will have undue insight when trying to predict the price in 2022. A general rule of thumb is that your model should not train on data that it would not have available when deployed. Therefore, you can use 2022 stock price data (in your train set) to predict 2023 values (in your test set), but not the other way around. These splitting strategies are specific to every dataset and ML problem. When designing a new pipeline with a non-IID dataset, think about how your model will be deployed in order to inform the splitting strategy you use to train and test it.

### 0.1.6 Problem 2a (10 points)

How would you split the dataset if we wanted to know how accurately we can predict the hand postures of a new, previously unseen user? What's the target variable? Write down your reasoning (the usual 1-2 paragraphs are fine). Split the dataset into training, validation, and test sets. As usual, check for reproducability!

Add your explanation here:

**Explanation:** From the original dataset, we can easily find that there are more than one data-point(row) for each user, so there may have a group structure problem, which means the dataset is non-IID. In order to deal with this problem, we need to split data by using GroupKFold method or GroupShuffleSplit method. Because there are 14 users which means 14 groups in dataset, it's better for us to use GroupShuffleSplit because the number of group is large.

Additionally, from the dataset itself and the data description, we could get that the target variable is 'Class', which represents the class ID of the given record, and it ranges from 1 to 5 with 1=Fist(with thumb out), 2=Stop(hand flat), 3=Point1(point with pointer finger), 4=Point2(point with pointer and middle fingers), 5=Grab(fingers curled as if to grab).

```
[ ]: # add your code here
posture_data = pd.read_csv("/Users/apple/Desktop/Data_1030/
                           ↪github-classroom-Data1030-Xiner Zhao/ps3-splitting-XXXXiner/data/Postures.
                           ↪csv")
print(posture_data.head())
```

	Class	User	X0	Y0	Z0	X1	Y1	\
0	0	0	0.000000	0.000000	0.000000	0.000000	0.000000	
1	1	0	54.263880	71.466776	-64.807709	76.895635	42.462500	
2	1	0	56.527558	72.266609	-61.935252	39.135978	82.538530	
3	1	0	55.849928	72.469064	-62.562788	37.988804	82.631347	
4	1	0	55.329647	71.707275	-63.688956	36.561863	81.868749	

	Z1	X2	Y2	...	Z8	X9	Y9	Z9	X10	Y10	Z10	X11	Y11	Z11
0	0.000000	0.000000	0.000000	...	0	0	0	0	0	0	0	0	0	0
1	-72.780545	36.621229	81.680557	...	?	?	?	?	?	?	?	?	?	?
2	-49.596509	79.223743	43.254091	...	?	?	?	?	?	?	?	?	?	?
3	-50.606259	78.451526	43.567403	...	?	?	?	?	?	?	?	?	?	?
4	-52.752784	86.320630	68.214645	...	?	?	?	?	?	?	?	?	?	?

[5 rows x 38 columns]

```
[ ]: df = posture_data.drop(labels=0, axis=0) # drop the first row
X = df.drop(columns=['Class', 'User'])
Y = df['Class']
groups = df['User']

print(X.head())
print(X.shape)
print(X.dtypes)

print(Y.head())
print(Y.shape)
print(Y.dtypes)
```

```
print(groups.head())
print(groups.shape)
print(groups.dtypes)
```

	X0	Y0	Z0	X1	Y1	Z1	\
1	54.263880	71.466776	-64.807709	76.895635	42.462500	-72.780545	
2	56.527558	72.266609	-61.935252	39.135978	82.538530	-49.596509	
3	55.849928	72.469064	-62.562788	37.988804	82.631347	-50.606259	
4	55.329647	71.707275	-63.688956	36.561863	81.868749	-52.752784	
5	55.142401	71.435607	-64.177303	36.175818	81.556874	-53.475747	

	X2	Y2	Z2	X3	...	Z8	X9	Y9	Z9	X10	Y10	\
1	36.621229	81.680557	-52.919272	85.2322638852917	...	?	?	?	?	?	?	
2	79.223743	43.254091	-69.982489	87.4508729469625	...	?	?	?	?	?	?	
3	78.451526	43.567403	-70.658489	86.8353875680762	...	?	?	?	?	?	?	
4	86.320630	68.214645	-72.228461	61.5961571288978	...	?	?	?	?	?	?	
5	76.986143	42.426849	-72.574743	86.3687480605765	...	?	?	?	?	?	?	

	Z10	X11	Y11	Z11
1	?	?	?	?
2	?	?	?	?
3	?	?	?	?
4	?	?	?	?
5	?	?	?	?

```
[5 rows x 36 columns]
(78095, 36)

X0      float64
Y0      float64
Z0      float64
X1      float64
Y1      float64
Z1      float64
X2      float64
Y2      float64
Z2      float64
X3      object
Y3      object
Z3      object
X4      object
Y4      object
Z4      object
X5      object
Y5      object
Z5      object
X6      object
Y6      object
Z6      object
```


```
X7      object
Y7      object
Z7      object
X8      object
Y8      object
Z8      object
X9      object
Y9      object
Z9      object
X10     object
Y10     object
Z10     object
X11     object
Y11     object
Z11     object
dtype: object
1      1
2      1
3      1
4      1
5      1
Name: Class, dtype: int64
(78095,)
int64
1      0
2      0
3      0
4      0
5      0
Name: User, dtype: int64
(78095,)
int64
```

```
[ ]: Y.value_counts()
```

```
[ ]: Class
3      16344
1      16265
5      15733
2      14978
4      14775
Name: count, dtype: int64
```

```
[ ]: plt.figure(figsize=(6,4))
pd.value_counts(Y).plot.barh()
plt.title('The distribution of Class')
```

```
[ ]: Text(0.5, 1.0, 'The distribution of Class')
```



```
[ ]: from sklearn.model_selection import GroupShuffleSplit

# Define the splitter with the desired sizes
gss = GroupShuffleSplit(n_splits=10, train_size=0.6, test_size=0.2, random_state=0)

for train_idx, test_idx in gss.split(X, Y, groups):
    # Split into train, validation, and test sets
    X_train = X.iloc[train_idx]
    Y_train = Y.iloc[train_idx]

    X_temp = X.iloc[test_idx]
    Y_temp = Y.iloc[test_idx]

    # Further split into validation and test sets
    val_size = 0.2 # Percentage of the data for validation
    val_split = int(val_size * len(X_temp))

    X_val = X_temp.iloc[:val_split]
    Y_val = Y_temp.iloc[:val_split]

    X_test = X_temp.iloc[val_split:]
    Y_test = Y_temp.iloc[val_split:]
```

```

# Print sizes for verification
print("Train:", X_train.index, '\n', "Validation:", X_val.index, '\n', u
↳"Test:", X_test.index)
print("Train size:", len(X_train))
print("Validation size:", len(X_val))
print("Test size:", len(X_test))
print('\n')

```

```

Train: Index([ 9049,  9050,  9051,  9052,  9053,  9054,  9055,  9056,  9057,
9058,
...
    78086, 78087, 78088, 78089, 78090, 78091, 78092, 78093, 78094, 78095],
dtype='int64', length=46414)
Validation: Index([18658, 18659, 18660, 18661, 18662, 18663, 18664, 18665,
18666, 18667,
...
    20551, 20552, 20553, 20554, 20555, 20556, 20557, 20558, 20559, 20560],
dtype='int64', length=1903)
Test: Index([20561, 20562, 20563, 20564, 20565, 20566, 20567, 20568, 20569,
20570,
...
    39353, 39354, 39355, 39356, 39357, 39358, 39359, 39360, 39361, 39362],
dtype='int64', length=7614)
Train size: 46414
Validation size: 1903
Test size: 7614

```

```

Train: Index([     1,      2,      3,      4,      5,      6,      7,      8,      9,
10,
...
    78086, 78087, 78088, 78089, 78090, 78091, 78092, 78093, 78094, 78095],
dtype='int64', length=48106)
Validation: Index([13766, 13767, 13768, 13769, 13770, 13771, 13772, 13773,
13774, 13775,
...
    15609, 15610, 15611, 15612, 15613, 15614, 15615, 15616, 15617, 15618],
dtype='int64', length=1853)
Test: Index([15619, 15620, 15621, 15622, 15623, 15624, 15625, 15626, 15627,
15628,
...
    28130, 28131, 28132, 28133, 28134, 28135, 28136, 28137, 28138, 28139],
dtype='int64', length=7416)
Train size: 48106
Validation size: 1853
Test size: 7416

```

```

Train: Index([ 1, 2, 3, 4, 5, 6, 7, 8, 9,
10,
...
78086, 78087, 78088, 78089, 78090, 78091, 78092, 78093, 78094, 78095],
dtype='int64', length=43104)
Validation: Index([ 9049, 9050, 9051, 9052, 9053, 9054, 9055, 9056,
9057, 9058,
...
12317, 12318, 12319, 12320, 12321, 12322, 12323, 12324, 12325, 12326],
dtype='int64', length=3278)
Test: Index([12327, 12328, 12329, 12330, 12331, 12332, 12333, 12334, 12335,
12336,
...
61852, 61853, 61854, 61855, 61856, 61857, 61858, 61859, 61860, 61861],
dtype='int64', length=13115)
Train size: 43104
Validation size: 3278
Test size: 13115

```

```

Train: Index([ 1, 2, 3, 4, 5, 6, 7, 8, 9,
10,
...
56987, 56988, 56989, 56990, 56991, 56992, 56993, 56994, 56995, 56996],
dtype='int64', length=51020)
Validation: Index([18658, 18659, 18660, 18661, 18662, 18663, 18664, 18665,
18666, 18667,
...
20740, 20741, 20742, 20743, 20744, 20745, 20746, 20747, 20748, 20749],
dtype='int64', length=2092)
Test: Index([20750, 20751, 20752, 20753, 20754, 20755, 20756, 20757, 20758,
20759,
...
61852, 61853, 61854, 61855, 61856, 61857, 61858, 61859, 61860, 61861],
dtype='int64', length=8370)
Train size: 51020
Validation size: 2092
Test size: 8370

```

```

Train: Index([ 1, 2, 3, 4, 5, 6, 7, 8, 9,
10,
...
70591, 70592, 70593, 70594, 70595, 70596, 70597, 70598, 70599, 70600],
dtype='int64', length=53540)
Validation: Index([18279, 18280, 18281, 18282, 18283, 18284, 18285, 18286,
18287, 18288,
...

```

```

    57263, 57264, 57265, 57266, 57267, 57268, 57269, 57270, 57271, 57272],
    dtype='int64', length=1147)
Test: Index([57273, 57274, 57275, 57276, 57277, 57278, 57279, 57280, 57281,
57282,
...
    61852, 61853, 61854, 61855, 61856, 61857, 61858, 61859, 61860, 61861],
    dtype='int64', length=4589)
Train size: 53540
Validation size: 1147
Test size: 4589

Train: Index([ 9049,  9050,  9051,  9052,  9053,  9054,  9055,  9056,  9057,
9058,
...
    70591, 70592, 70593, 70594, 70595, 70596, 70597, 70598, 70599, 70600],
    dtype='int64', length=43974)
Validation: Index([    1,     2,     3,     4,     5,     6,     7,     8,     9,    10,
...
    2920, 2921, 2922, 2923, 2924, 2925, 2926, 2927, 2928, 2929],
    dtype='int64', length=2929)
Test: Index([ 2930,  2931,  2932,  2933,  2934,  2935,  2936,  2937,  2938,
2939,
...
    28622, 28623, 28624, 28625, 28626, 28627, 28628, 28629, 28630, 28631],
    dtype='int64', length=11716)
Train size: 43974
Validation size: 2929
Test size: 11716

Train: Index([18658, 18659, 18660, 18661, 18662, 18663, 18664, 18665, 18666,
18667,
...
    78086, 78087, 78088, 78089, 78090, 78091, 78092, 78093, 78094, 78095],
    dtype='int64', length=50196)
Validation: Index([ 9049,  9050,  9051,  9052,  9053,  9054,  9055,  9056,
9057,  9058,
...
    11858, 11859, 11860, 11861, 11862, 11863, 11864, 11865, 11866, 11867],
    dtype='int64', length=2819)
Test: Index([11868, 11869, 11870, 11871, 11872, 11873, 11874, 11875, 11876,
11877,
...
    61852, 61853, 61854, 61855, 61856, 61857, 61858, 61859, 61860, 61861],
    dtype='int64', length=11276)
Train size: 50196
Validation size: 2819

```

```

Test size: 11276

Train: Index([    1,      2,      3,      4,      5,      6,      7,      8,      9,
10,
...
78086, 78087, 78088, 78089, 78090, 78091, 78092, 78093, 78094, 78095],
dtype='int64', length=54947)
Validation: Index([13766, 13767, 13768, 13769, 13770, 13771, 13772, 13773,
13774, 13775,
...
16504, 16505, 16506, 16507, 16508, 16509, 16510, 16511, 16512, 16513],
dtype='int64', length=2748)
Test: Index([16514, 16515, 16516, 16517, 16518, 16519, 16520, 16521, 16522,
16523,
...
70591, 70592, 70593, 70594, 70595, 70596, 70597, 70598, 70599, 70600],
dtype='int64', length=10996)
Train size: 54947
Validation size: 2748
Test size: 10996

Train: Index([13766, 13767, 13768, 13769, 13770, 13771, 13772, 13773, 13774,
13775,
...
78086, 78087, 78088, 78089, 78090, 78091, 78092, 78093, 78094, 78095],
dtype='int64', length=47785)
Validation: Index([28140, 28141, 28142, 28143, 28144, 28145, 28146, 28147,
28148, 28149,
...
30563, 30564, 30565, 30566, 30567, 30568, 30569, 30570, 30571, 30572],
dtype='int64', length=2433)
Test: Index([30573, 30574, 30575, 30576, 30577, 30578, 30579, 30580, 30581,
30582,
...
61852, 61853, 61854, 61855, 61856, 61857, 61858, 61859, 61860, 61861],
dtype='int64', length=9735)
Train size: 47785
Validation size: 2433
Test size: 9735

Train: Index([    1,      2,      3,      4,      5,      6,      7,      8,      9,
10,
...
78086, 78087, 78088, 78089, 78090, 78091, 78092, 78093, 78094, 78095],
dtype='int64', length=36614)

```

```

Validation: Index([35443, 35444, 35445, 35446, 35447, 35448, 35449, 35450,
35451, 35452,
...
49150, 49151, 49152, 49153, 49154, 49155, 49156, 49157, 49158, 49159],
dtype='int64', length=4144)
Test: Index([49160, 49161, 49162, 49163, 49164, 49165, 49166, 49167, 49168,
49169,
...
70591, 70592, 70593, 70594, 70595, 70596, 70597, 70598, 70599, 70600],
dtype='int64', length=16576)
Train size: 36614
Validation size: 4144
Test size: 16576

```

**Reproducibility Check:** We can rerun the code cell above manually for many times, and we can see that the index of training set, validation set, test set is always same no matter how many times I rerun the GroupShuffleSplit code above.

### 0.1.7 Problem 2b (10 points)

How would you split the data if we wanted to identify a user based on their hand postures? What's the target variable? Follow the same steps as in 2a (explain your reasoning, split, check reproducability).

Add your explanation here:

**Explanation:** Because we do not need to predict the hand postures from a new, previously unseen user, but we need to identify a user based on their hand postures, which means we need to identify out which users that hand postures belong to based on available information. So, the target variable should be 'User', but not 'Class'. Also, because different users have different numbers of datapoints, which means some of the users have more datapoint but some of them have less, or even their dataset are 'rare'. Therefore, in order to use all the information of hand postures from all users, we need to split the datapoint of each user into train, validation, and test data set evenly. To achieve that, we need to choose StratifiedKFold splitting method.

```
[ ]: groups.value_counts()
```


```
[ ]: User
```

10	9573
0	9048
13	8739
11	8061
14	7495
8	6811
5	5105
12	4865

```
1    4717  
2    4513  
6    4377  
9    3920  
7    492  
4    379  
Name: count, dtype: int64
```

```
[ ]: plt.figure(figsize = (6,4))  
pd.value_counts(groups).plot.bart()  
plt.title('The quantity of datapoint of each user')
```

```
[ ]: Text(0.5, 1.0, 'The quantity of datapoint of each user')
```



```
[ ]: # add your code here  
from sklearn.model_selection import train_test_split  
from sklearn.model_selection import StratifiedKFold  
  
random_state = 1030  
# stratified train-test split  
X_other, X_test, Y_other, Y_test = train_test_split(X, groups, test_size = 0.  
        ↵2, stratify=groups, random_state=random_state)  
print('test balance:', np.unique(Y_test, return_counts=True))
```

```

# do StratifiedKFold split on other
kf = StratifiedKFold(n_splits=4,shuffle=True,random_state=random_state)
for train_index, val_index in kf.split(X_other,Y_other):
    print('new fold')
    X_train = X_other.iloc[train_index]
    y_train = Y_other.iloc[train_index]
    X_val = X_other.iloc[val_index]
    y_val = Y_other.iloc[val_index]
    print(np.unique(y_train,return_counts=True))
    print(np.unique(y_val,return_counts=True))

test balance: (array([ 0,  1,  2,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]),
array([1810,  943,  903,   76, 1021,  875,   98, 1362,  784, 1915, 1612,
       973, 1748, 1499]))
new fold
(array([ 0,  1,  2,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]), array([5428,
2830, 2707, 228, 3063, 2626, 296, 4087, 2352, 5744, 4836,
2919, 5244, 4497]))
(array([ 0,  1,  2,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]), array([1810,
944,  903,   75, 1021,  876,   98, 1362,  784, 1914, 1613,
       973, 1747, 1499]))
new fold
(array([ 0,  1,  2,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]), array([5428,
2830, 2708, 227, 3063, 2627, 295, 4087, 2352, 5744, 4837,
2919, 5243, 4497]))
(array([ 0,  1,  2,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]), array([1810,
944,  902,   76, 1021,  875,   99, 1362,  784, 1914, 1612,
       973, 1748, 1499]))
new fold
(array([ 0,  1,  2,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]), array([5429,
2831, 2708, 227, 3063, 2627, 295, 4086, 2352, 5743, 4837,
2919, 5243, 4497]))
(array([ 0,  1,  2,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]), array([1809,
943,  902,   76, 1021,  875,   99, 1363,  784, 1915, 1612,
       973, 1748, 1499]))
new fold
(array([ 0,  1,  2,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]), array([5429,
2831, 2707, 227, 3063, 2626, 296, 4087, 2352, 5743, 4837,
2919, 5243, 4497]))
(array([ 0,  1,  2,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14]), array([1809,
943,  903,   76, 1021,  876,   98, 1362,  784, 1915, 1612,
       973, 1748, 1499]))

```

**Reproducibility Check:** We can rerun the code cell above manually for many times, and we can see that the index of training set, validation set, test set is always same no matter how many times I rerun the StratifiedKFold code above.